

Is a Plan Load-Bearing? A Black-Box Causal Audit of Plan-Based Oversight for LLM Agents

Kareem Fareed
Stanford University
kfareed@stanford.edu

Abstract

Plan and CoT-based monitoring of LLM agents assumes that the plan an agent states is behaviorally meaningful, i.e., that reading the plan tells an overseer which irreversible actions the model will take. This assumption has been highlighted as unvalidated by recent monitorability agendas, and no black-box experimental design tests it over tool-executing trajectories. We introduce a validity-gated counterfactual test: on τ^2 -bench tasks rewritten so that a customer is explicitly indifferent between two outcomes (e.g., refund vs. exchange), we verify behaviorally that both outcomes are achievable and policy-compliant, have the agent ratify a plan for each, then counterfactually seed one or the other as the agent’s own stated plan and measure which tools the agent calls. Across four open-weight models on the same 62 tasks, the seeded plan causally shifts the outcome in every model. Binding ranges from 0.75 to 0.93; for Qwen3-30B, the probability of outcome A moves from 0.97 under plan A to 0.11 under plan B. Failures correlate with the model’s *revealed preference*, measured in separate unseeded rollouts. In all four models, overrides are rare when the seeded plan agrees with this preference (0–8%) and multiply when it conflicts (12–49%); task-clustered bootstrap intervals for the conflict-minus-agree difference exclude zero in every model. Overrides are silent, uncorrelated with the model’s own stated validity judgments, and model-dependent. Between 5 and 7 tasks per model are fully seed-invariant, and all 19 such tasks with a measurable preference realize the preferred outcome. These results support plan-based oversight as conditionally valid, and yield a practical black-box audit: measure an agent’s preferences with a few unseeded rollouts, and do not trust stated plans that conflict with them.

1 Introduction

A growing share of AI oversight practice relies on reading what agents *say they will do*. Chain-of-thought and plan monitoring are attractive because they are cheap, black-box, and often empirically informative (Korbak et al., 2025; Guan et al., 2025). But their safety value rests on the premise that the visible plan actually constrains the action output. If an agent’s stated plan is decorative, and it would execute the same actions regardless of what the plan says, then a monitor reading plans learns nothing about what the agent is going to do.

This premise is, at present, an assumption rather than a measured property. The groups building monitorability evaluations name its validation as open work (Guan et al., 2025; Korbak et al., 2025), and evaluation methodologists caution that causal claims about agent behavior require experimental or pseudo-experimental designs that trajectory analysis currently lacks (Kirgis et al., 2026). Since plan-reading is being adopted as an oversight layer for agents that take irreversible actions on real systems, there are real stakes in knowing if plans can silently fail to govern an agent’s action.

Testing this premise isn’t as simple as perturbing the agent’s plan and observing its behavior. If the substituted plan is worse than the agent’s own, ignoring it is good judgment, not unfaithfulness. Moreover, if the task has only one valid outcome, following any plan is

indistinguishable from following none. What is needed is a black-box intervention that swaps between two *genuinely valid, agent-endorsed* plans whose outcomes differ in a way an overseer cares about, with the agent’s subsequent tool calls as the outcome measure. Existing interventions remove plans, corrupt reasoning, or impose external instructions; none meet this bar (§2). We contribute such a design and its first results.

The diagnostic. We construct tasks with two genuinely interchangeable outcomes: τ^2 -bench (Barres et al., 2025) customer-service episodes rewritten so the customer is explicitly, symmetrically indifferent between two resolutions (money back vs. replacement item; canceling vs. adjusting a pending order), with all factual details preserved. Fixed template plans, one per outcome, are *ratified* by the model, then *force-executed* to verify behaviorally that each outcome is achievable and policy-compliant in the environment. On tasks passing these gates, we run counterbalanced free rollouts with one plan seeded as the agent’s own stated plan, and score which outcome’s write action actually succeeds in the environment. Separately, *unseeded* rollouts (no plan at all) measure the model’s revealed preference between the outcomes.

Findings. Across 62 validity-gated tasks and four open-weight models on identical inputs we find:

1. **Endorsed plans causally shift actions in every model, to varying degrees.** Seeding plan A vs. plan B moves $P(\text{outcome A})$ from 0.97 to 0.11 for the strongest binder, and overall binding spans 0.75–0.93 (Table 1), exceeding each model’s same-seed variability floor of 4.5–15.3%. This is direct evidence for the premise of plan-based monitoring.
2. **Failures are directional and model-specific.** Qwen3-30B’s overrides favor outcome A (13:3) and gpt-oss-120b’s favor outcome B (3:57), while the GLM-family models are less globally directional (20:11 and 20:18). On 5–7 of 62 tasks per model, the agent produces the same scored outcome in every rollout regardless of which plan it stated.
3. **Revealed preference predicts where the plan fails; stated judgment does not.** Overrides are rare when the seeded plan agrees with the model’s separately measured unseeded preference and multiply when it conflicts, in all four models (Table 1). The task-clustered contrast excludes zero in every model, and all 19 seed-invariant tasks with a measurable preference realize exactly the preferred outcome. Ratification judgments, by contrast, provide no positive signal (§5.2).
4. **Plans play different functional roles across models.** Without any plan, gpt-oss-120b fails to complete a successful write in 82% of rollouts, while Qwen3-30B fails in 12%. For one model the plan is close to an actuator, whereas for the other it is a suggestion weighed against the model’s preferences.

Implication. Plan-based oversight is neither invalid nor fully reliable. It is *conditionally causal*, and the condition is measurable from a handful of black-box samples, which yields the audit we develop in §6.

We release the harness, regenerated task rewrites, and rollout-level data for all four models.¹

2 Related Work

CoT monitorability. Korbak et al. (2025) pose the evaluation of CoT monitoring as open, and Guan et al. (2025) name validating it for load-bearing use as future work, describing “intervention” archetypes of exactly the form we instantiate. We differ in operating on tool-executing agents and measuring causal control of *actions* rather than monitor detection rates.

¹<https://github.com/KookiesNKareem/plan-load-bearing>. The raw logs and rewrite cache of the original 59-task runs were lost in a storage failure after analysis. Every quantitative result in this version is recomputed from the released independent regeneration on its 62-task canonical set.

Interventions on reasoning. Causal tests that corrupt or ablate reasoning chains find mixed faithfulness in QA settings (Khanzadeh, 2026; Han et al., 2026), but corruption cannot distinguish unfaithfulness from justified correction, since ignoring a broken chain is good behavior. Our intervention swaps between two *validated, ratified* alternatives, removing that excuse. Mehta & Datta (2026) study plan *removal*: white-box probes plus an eviction stress-test showing that deleting the plan degrades task success. Removal establishes that plans matter, but cannot reveal which of two valid outcomes a stated plan selects, the counterfactual our swap isolates.

Plan adherence. Liu et al. (2026) measure whether programming agents follow externally supplied plans. Externally imposed plans confound compliance with judgment. Our plans are ratified by the model as valid resolutions before seeding, and both alternatives are verified achievable in the environment.

Agent evaluation methodology. Kirgis et al. (2026) call for experimental designs over agent trajectories, and our paradigm is one. We build on τ -bench and τ^2 -bench (Yao et al., 2024; Barres et al., 2025), whose reference-scored tasks pin a unique resolution, which is why our diagnostic must rewrite requests to open a second valid outcome (§3).

3 A Validity-Gated Counterfactual Plan-Seeding Diagnostic

Setting. We use the τ^2 -bench retail environment: an agent with 14 tools (lookups plus state-changing “writes” such as `exchange_delivered_order_items`, `return_delivered_order_items`, `modify_pending_order_items`, `cancel_pending_order`) operating under a policy document against a database of users, orders, and products. Tasks whose reference solutions touch delivered or pending orders admit a canonical *outcome pair*: exchange vs. return for delivered orders, modify vs. cancel for pending ones. The two outcomes differ in which irreversible write occurs, exactly the distinction an overseer of consequential actions cares about.

Opening the second outcome. Benchmark tasks dictate the outcome (“you want to exchange...”), as they must to be scorable against a unique reference. We therefore rewrite each task’s request so the customer states the underlying problem and is *explicitly, symmetrically indifferent*. Money back would fully satisfy them (stated first), and a suitable replacement would work equally well. Every factual detail (order IDs, items, attribute preferences as a conditional aside) is preserved verbatim, and nothing absent from the original may be invented. Rewrites are generated once, screened for outcome-verb leakage, and shared across all models. We view this construction the way psychophysics views a staircase procedure: not as deployment realism, but as the controlled regime in which the causal question is answerable. If plan-reading fails even here, deployment claims inherit the doubt.

Validity gates. For each task we instantiate two fixed template plans (shared lookups + the outcome’s write). Three gates must pass:

1. **Ratification.** The model judges, given the policy and rewritten request, whether each plan’s end state resolves the request compliantly. This preserves self-endorsement (overriding a plan the model itself ratified cannot be excused as correcting a bad plan) without the brittleness of free-form authoring, and is recorded as a covariate rather than used to censor.
2. **Forced execution.** Each plan is executed under instruction in a fresh environment (lookups free, the designated write mandatory). The task passes only if *both* writes succeed against the live environment, behavioral proof that both outcomes are achievable rather than merely plausible.
3. **Success-only scoring.** In all analysis, an outcome counts only if its write *succeeded* in the environment. Scoring attempted-but-failed writes produced a spurious perfect binding rate in piloting.

Measurement. On each gated task we run: (i) four counterbalanced *seeded* rollouts, in which the plan (A twice, B twice) is placed in the transcript as the agent’s own stated plan (“My

Model	scored rollouts	$P(\text{outcome A} \mid \text{seed})$		binding	seed- invariant	$P(\text{override} \mid \text{pref.})$	
		plan A	plan B			agrees	conflicts
Qwen3-30B-A3B	229/248	0.972	0.107	0.930	5/62	0.000	0.122
gpt-oss-120b	241/496	0.467	0.022	0.751	6/62	0.023	0.488
GLM-4.5-Air	131/248	0.788	0.253	0.763	6/62	0.081	0.415
GLM-4.7-Flash	212/248	0.818	0.177	0.821	7/62	0.032	0.341

Table 1: Per-model summary on the 62 regenerated canonical tasks (identical inputs). *Scored* counts rollouts with an environment-successful write; *binding* is the rate at which the realized outcome matches the seeded plan; *seed-invariant* counts tasks whose outcome never varies with the seeded plan; the last two columns give override rates when the seeded plan agrees vs. conflicts with the model’s separately sampled unseeded revealed preference. gpt-oss-120b pools two identical passes (eight seeded and eight unseeded rollouts per task; its seed-invariance requires all eight seeded rollouts to agree). Wilson CIs in text and Figure 2.

plan: ...I will execute it now”) and the agent then acts freely; (ii) four *unseeded* rollouts with no plan, whose realized outcomes define the model’s revealed preference for the task. Same-seed pairs give a per-task nondeterminism floor, and the forced runs double as a positive control. An LLM judge scores every rollout for resolution and policy compliance, with all judgments re-issued by a single fixed judge for the robustness analysis. Figure 1 summarizes the pipeline.

Metrics. For a rollout seeded with plan X, define *binding* as the event that outcome X’s write is the one that succeeds, and *override* as the complementary event where the opposite outcome’s write succeeds. Task-level *seed-invariance*, the same outcome realized in all four seeded rollouts with zero variance, is the signature of a plan with no causal effect.

4 Experimental Setup

We evaluate **Qwen3-30B-A3B** (official FP8) (Qwen Team, 2025), **gpt-oss-120b** (native MXFP4) (OpenAI, 2025), **GLM-4.5-Air** (community 4-bit AWQ, the one non-official quantization in our set, a caveat we carry into §7), and **GLM-4.7-Flash** (official BF16 weights; the newest model in our set, released January 2026), served locally with vLLM (Kwon et al., 2023) at temperature 0 throughout (server-side batching still yields measurable nondeterminism floors, which we report rather than hide). Of 114 retail tasks, 96 admit an outcome pair, and 62 pass all gates under Qwen3-30B to form the regenerated canonical task set. All four models are evaluated on this same set with the same cached rewrites and eight concurrent task workers. gpt-oss-120b was run twice through the identical protocol, so its rollout counts are doubled throughout (496 rather than 248; one task in its second seeded pass failed before producing rollouts). For the other three models the gates are recorded as covariates rather than used to censor, since model-specific gating (for example, gpt-oss-120b’s stricter ratification) would otherwise select a different, biased subset per model. A second-domain replication (τ^2 -bench airline; update-reservation vs. cancel-and-rebook as write-sets) probes whether the protocol transfers (§5.5).

5 Results

5.1 Endorsed plans causally control actions, mostly

Seeding plan A versus plan B moves the probability of realizing outcome A from 0.972 [0.921, 0.991] to 0.107 [0.064, 0.175] for Qwen3-30B, from 0.467 [0.376, 0.561] to 0.022 [0.008, 0.064] for gpt-oss-120b, from 0.788 [0.660, 0.878] to 0.253 [0.170, 0.359] for GLM-4.5-Air, and from 0.818 [0.731, 0.882] to 0.177 [0.118, 0.258] for GLM-4.7-Flash (Figure 2, left; Table 1). Overall binding is 0.930, 0.751, 0.763, and 0.821, respectively. Same-seed nondeterminism floors are 4.5% for Qwen, 4.9% for gpt-oss, 11.5% for GLM-4.5, and 15.3% for GLM-4.7, below the plan-induced shift in every model. Thus an overseer reading the stated plan predicts




Figure 1: The diagnostic. A template plan per outcome is ratified and force-validated, then seeded as the agent’s own stated plan in counterbalanced free rollouts scored by which irreversible write succeeds. Unseeded rollouts measure revealed preference.




Figure 2: **Left:** the seeded plan shifts the realized outcome in all four models (Wilson 95% CIs over scored rollouts). **Right:** overrides concentrate under preference conflict. Per-rollout Wilson CIs are descriptive; task-clustered intervals are reported in §5.3. Rollouts per bar (agree/conflict): Qwen 105/98, gpt-oss 44/41 (pooled over two passes), GLM-4.5 62/53, and GLM-4.7 95/91.

the successful irreversible action in roughly three quarters to nine tenths of scored rollouts. Filtering to transcripts that a single fixed judge marks resolved and policy-compliant retains binding of 0.937 (89/95), 0.719 (92/128), 0.750 (57/76), and 0.800 (96/120); because resolution is post-treatment, these are robustness checks rather than primary estimands. Because rollouts cluster within tasks, per-rollout Wilson intervals are descriptive and the preference-conflict contrasts below use task-level resampling (§5.3).

5.2 Failures are directional, model-specific, and sometimes total

The rollouts that escape the plan do not scatter. Qwen3-30B’s 16 overrides favor outcome A (13:3), while gpt-oss-120b’s 60 overrides overwhelmingly favor outcome B (3:57). The GLM-family runs are less globally directional (20:11 for GLM-4.5-Air and 20:18 for GLM-4.7-Flash). The unseeded data sharpen the distinction: global $P(A)$ is 0.530, 0.132, 0.404, and 0.523, respectively, while the numbers of measurable preference-A versus preference-B tasks are 28:25, 1:15, 16:27, and 27:27. Thus the robust result is model-specific preference structure, not the submitted claim that Qwen and gpt-oss carry mirror-image global preferences. At the

extreme, 5 Qwen tasks, 6 gpt-oss tasks (requiring agreement across all eight seeded rollouts), 6 GLM-4.5 tasks, and 7 GLM-4.7 tasks are *seed-invariant*: every scored rollout realizes the same outcome regardless of which plan the agent stated. The invariant outcomes split A:B as 5:0, 0:6, 3:3, and 3:4. On all 19 invariant tasks with a measurable unseeded preference, the outcome matches that preference (§5.3).

The GLM pair also gives a within-family comparison: GLM-4.7-Flash (official BF16, one generation newer) binds at 0.821 against its sibling’s 0.763, loses fewer conflicts (0.341 vs. 0.415), and completes writes more reliably (85% vs. 53% of seeded rollouts scored), while both retain task-specific preferences. Generation and precision change together here, so their contributions cannot be separated.

Ratification does not *predict* behavior in any model where refusals permit the comparison (Qwen’s gated run retains only ratified pairs). Per-rollout override rates on ratified-both versus refused-one tasks are 0.209 vs. 0.286 for gpt-oss-120b, 0.250 vs. 0.205 for GLM-4.5-Air, and 0.143 vs. 0.192 for GLM-4.7-Flash. The small point-estimate differences change direction across models. In our data, what a model *says* about plan validity provides no positive signal about whether it will follow the plan.

5.3 Revealed preference predicts where plans fail, and the prediction survives held-out splits

Unseeded rollouts reveal stable per-task preferences between the two outcomes. Each task’s preference is measured from its own $K=4$ unseeded rollouts ($K=8$ for pooled gpt-oss). Modal labels agree between independent sample halves on 46/46 comparable Qwen tasks, 6/6 gpt-oss tasks across its two passes, 25/25 GLM-4.5 tasks, and 26/27 GLM-4.7 tasks. Aggregate $P(A)$ differs sharply across models (0.530, 0.132, 0.404, and 0.523), so only gpt-oss has a strong global direction; the other models’ useful signal is primarily task-specific. Because temperature-0 decoding could make stability an artifact of determinism, we repeat the measurement with six rollouts per task at temperature 0.7. For Qwen, the modal label agrees with temperature 0 on 41/47 comparable tasks, with mean modal share 0.901; the temperature-0.7 labels still predict 2/100 agree-condition versus 13/91 conflict-condition overrides (0.020 vs. 0.143). For GLM-4.7, agreement is 46/50 and mean modal share is 0.840; its temperature-0.7 labels likewise preserve the contrast (5/98 vs. 30/99, or 0.051 vs. 0.303). Preference is therefore softer under sampling but remains behaviorally predictive.

This preference predicts *where* plan-binding fails (Figure 2, right). Tasks need at least two successful unseeded rollouts and an untied modal preference to enter this analysis. Qwen3-30B overrides 0 of 105 preference-agreeing plans versus 12 of 98 conflicts (0.000 [0.000, 0.035] vs. 0.122 [0.071, 0.202]); the task-clustered 95% interval for conflict minus agree is [0.043, 0.212]. The corresponding contrasts are 1/44 versus 20/41 for gpt-oss (0.023 vs. 0.488; clustered difference CI [0.175, 0.763]), 5/62 versus 22/53 for GLM-4.5 (0.081 vs. 0.415; [0.132, 0.509]), and 3/95 versus 31/91 for GLM-4.7 (0.032 vs. 0.341; [0.168, 0.420]). Thus the preference-conflict effect is positive under task-level resampling in every model. The agree-condition rate also sits below each model’s same-seed nondeterminism floor, consistent with ordinary serving variability carrying the residual agree-condition failures. Pooled as a descriptive summary, overrides occur in 9/306 agreeing rollouts (2.9% [1.6, 5.5]) versus 85/283 conflicts (30.0% [25.0, 35.6]). All 19 seed-invariant tasks with a measurable preference realize the preferred outcome (4/4 Qwen, 2/2 gpt-oss, 6/6 GLM-4.5, and 7/7 GLM-4.7). Plans nevertheless retain causal force throughout: even under conflict, the seeded plan wins a majority of scored rollouts in every model.

5.4 Plan-as-actuator vs. plan-as-suggestion

The models relate to plans differently at a functional level. Without any seeded plan, gpt-oss-120b fails to complete a successful write in 82% of rollouts (and 51% of its seeded rollouts are unscorable), versus 12%/8% for Qwen3-30B, 13%/15% for GLM-4.7-Flash, and 31%/47% for GLM-4.5-Air. Because each gpt-oss pass yields few scored rollouts, it runs two identical passes that we pool throughout. For gpt-oss, a plan is close to an *actuator*: the model frequently does not complete a write without one, and once acting under its

strongly preferred outcome-B plan it follows it in 97.8% of scored rollouts. For Qwen, a plan is a *suggestion*: action happens with or without it, and the plan competes with task-level preferences. Both regimes matter for oversight. All gpt-oss quantitative claims are conditional on the completed-write stratum, a selection we flag explicitly.

5.5 Second domain: the protocol transfers; power does not

On τ^2 -bench airline we instantiate the pair update-reservation vs. cancel-and-rebook, scored over write-*sets* rather than single writes, with mechanism-indifferent rewrites (the goal state is fixed; the two bookkeeping paths differ in which irreversible writes fire). Of 13 flight-change tasks, 8 fail the forced-execution gate because the 30B agent cannot reliably complete flight search and multi-argument booking even under instruction, a domain-competence floor we report rather than patch. On the 5 gated tasks, 17 of 18 scored seeded rollouts bind (0.944), including 5 of 6 whose seeded plan conflicts with a measurable unseeded preference. We present this replication as evidence that the paradigm transfers across domains and action-space geometries, not as independent confirmation of the conflict-override relationship.

6 Implications for Plan-Based Oversight

The premise holds, conditionally. These results are, to our knowledge, the first black-box causal evidence that an endorsed plan controls a tool-executing agent’s irreversible actions, which is the assumption plan-based monitoring rests on. An overseer reading plans is not reading theater.

But the failure mode is adversarially shaped. The failure channel (7–25% of scored rollouts, depending on model) is silent (no verbal signal accompanies an override), directional (it substitutes the model’s preferred irreversible action), model-specific, and, critically, invisible to the model’s own stated judgments. A monitor that asked the agent to double-check its plan’s validity would, on our evidence, learn nothing, because ratification and behavior are decoupled. The plan is least informative precisely when the agent has a live preference about the outcome, which is when oversight matters most. One qualification is important. Our conflicts are manufactured, whereas in deployment a self-authored plan and the subsequent action share the preference as a common cause and will usually agree, so the base rate of silent override in the wild may be much lower than the conflict-condition rate here. What we establish is the failure mode’s existence, direction, and predictability, and the audit below lets an overseer bound its base rate per decision class.

A practical audit. Because revealed preference is stable and cheap to measure, the diagnostic inverts into a deployment-time procedure. For a consequential decision class, sample the agent’s behavior K times without a plan and estimate its preference distribution. Stated plans that conflict with a concentrated preference are unreliable oversight surfaces (a $\sim 10\times$ override risk in our data) and should be routed to stronger oversight. The audit requires no internals, no training, and a handful of rollouts.

For monitorability research. Our design instantiates the “intervention” archetype of Guan et al. (2025) at the action-trajectory level requested by Kirgis et al. (2026). The validity gates are the load-bearing components. Without them, piloting produced large spurious effects in both directions (fake perfect binding from attempted-write scoring, fake theater from unvalidated alternatives).

7 Limitations

Ratified templates, not authored plans. We test whether a plan the model *endorses* controls its actions once stated as its own, not whether naturally generated plans are faithful descriptions of internal deliberation. It is the right primitive for monitor validity (the monitor reads whatever plan appears), with self-authored plans the natural extension. Ratification itself may be shallow agreement, but the causal effect of the seeded plan is large regardless of

endorsement depth. **Constructed indifference.** The rewritten requests create a decision regime where latent preference has nowhere to hide except action choice. This is a validity probe rather than an ecology, and deployment-distribution effect sizes may differ. The rewrites state the money-back option first; the regenerated Qwen run has no strong global outcome bias ($P(A) = 0.53$), so rewrite order effects remain uncontrolled rather than ruled out. **Scale and models.** Four open-weight models, one primary domain, $n=62$ tasks, and a power-limited second domain ($n=5$ gated tasks). Frontier API models remain untested. GLM-4.5-Air runs a community 4-bit quantization; GLM-4.7-Flash is full precision and newer, so its stronger binding (0.821 vs. 0.763) cannot be attributed separately to generation or precision. **Measurement.** Conflict labels derive from four same-task unseeded samples (eight for gpt-oss), so some modal labels remain uncertain even though the seeded and unseeded samples are independent; agreement between sample halves is high but not perfect. Primary runs decode at temperature 0, so preference stability could partly reflect decoding determinism; §5.3 reports temperature-0.7 replications. **Data availability.** The raw logs and rewrite cache of the original 59-task runs were lost after analysis. We therefore do not use those unreleasable values in this version: all tables, figures, and quantitative claims are recomputed from the released 62-task regeneration for all four models.

8 Conclusion

We turned an assumption of plan-based oversight into a measurable, black-box quantity, and found it conditionally true. Endorsed plans causally control agent actions except where they collide with a stable, model-specific revealed preference, and the collision is silent, self-report-invisible, and predictable from a few unplanned samples. We offer the paradigm, the confound battery, and the audit as infrastructure for the monitorability agenda.

References

- Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. τ^2 -bench: Evaluating conversational agents in a dual-control environment. *arXiv preprint arXiv:2506.07982*, 2025.
- Melody Y. Guan, Miles Wang, Micah Carroll, Zehao Dou, Annie Y. Wei, Marcus Williams, Benjamin Arnav, Joost Huizinga, Ian Kivlichan, Mia Glaese, Jakub Pachocki, and Bowen Baker. Monitoring monitorability. *arXiv preprint arXiv:2512.18311*, 2025.
- Yunseok Han, Yejoon Lee, and Jaeyoung Do. RFEval: Benchmarking reasoning faithfulness under counterfactual reasoning intervention in large reasoning models. In *International Conference on Learning Representations (ICLR)*, 2026.
- Sourena Khanzadeh. Project ariadne: A structural causal framework for auditing faithfulness in LLM agents. *arXiv preprint arXiv:2601.02314*, 2026.
- Peter Kirgis, Sayash Kapoor, Stephan Rabanser, Nitya Nadgir, Cozmin Ududec, Magda Dubois, JJ Allaire, Conrad Stosz, Marius Hobbhahn, Jacob Steinhardt, and Arvind Narayanan. Log analysis is necessary for credible evaluation of AI agents. *arXiv preprint arXiv:2605.08545*, 2026.
- Tomek Korbak, Mikita Balesni, Elizabeth Barnes, Yoshua Bengio, Joe Benton, Joseph Bloom, et al. Chain of thought monitorability: A new and fragile opportunity for AI safety. *arXiv preprint arXiv:2507.11473*, 2025.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, 2023.
- Shuyang Liu, Saman Dehghan, Jatin Ganhotra, Martin Hirzel, and Reyhaneh Jabbarvand. Evaluating plan compliance in autonomous programming agents. *arXiv preprint arXiv:2604.12147*, 2026.

Aman Mehta and Anupam Datta. Plans don't persist: Why context management is load bearing for LLM agents. *arXiv preprint arXiv:2606.22953*, 2026.

OpenAI. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025.

Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*, 2024.

A Reproducibility details

Serving. All models served with vLLM 0.24 on a single 96GB GPU; FP8 (Qwen3-30B-A3B-Instruct-2507), MXFP4 (gpt-oss-120b, reasoning effort low), 4-bit AWQ (GLM-4.5-Air, thinking mode), and BF16 (GLM-4.7-Flash, thinking mode); temperature 0; structured outputs (JSON schema) for all ratification and judging calls, which eliminated a schema-drift failure mode that silently discarded roughly half of free-form elicitations in piloting.

Funnel. Retail: 114 tasks $\rightarrow$ 96 with an outcome pair $\rightarrow$ 62 canonical (Qwen gates: 17 not ratified, 17 not force-executable). Per task: 2 forced + 4 seeded + 4 unseeded rollouts (doubled for gpt-oss-120b via a second identical pass), each rollout ≤ 16 tool calls with a bounded continue-nudge on non-tool turns. All primary runs use eight concurrent task workers; maximum generation is 600 tokens for Qwen, 2,000 for gpt-oss, and 2,500 for both GLM models.

Judging. Every seeded rollout is judged for resolution and policy compliance. The robustness analysis re-judges all 1,236 regenerated free-rollout transcripts with a single fixed judge (gpt-oss-120b) given the pre-confirmation context (the seeded transcripts stipulate that the customer has already confirmed, so a missing confirmation step is not a violation).

Known pitfalls (each cost a pilot iteration and is controlled in the released harness): scoring attempted rather than environment-successful writes; outcome-pinned benchmark requests; rewriters fabricating absolute details (order numbers, flight numbers) when the source request is relative; exact-sequence forcing conflating recipe-following failures with outcome invalidity; per-model gating silently selecting non-comparable task subsets across models.

B Worked example: a seed-invariant task

All text below comes from the released regeneration (retail task 45, one of the five Qwen3-30B seed-invariant tasks).

Original benchmark request (outcome-pinned):

"You want to exchange a robotic vacuum cleaner in your recent order for a canister based one from the same product line. When asked for order ID, provide 9502127 first. If that doesn't work, respond exactly with 'I forgot the W at the beginning'. If and only if the agent gives you several options for the new vacuum, go for the bagless version." (*Known info: Daiki Johnson, Denver, zip 80273.*)

Indifference rewrite (cached; identical input for all models):

"I don't need the robotic vacuum cleaner from my recent order, and simply getting my money back would fully satisfy me. If I do end up with a different item, it should be exactly: a canister-style vacuum from the same product line, bagless. Whichever you recommend is fine—I truly have no preference."

Template plans (both ratified by the model; both force-validated in the environment):

Plan A: find_user_id_by_name_zip $\rightarrow$ get_user_details $\rightarrow$ get_order_details $\rightarrow$ get_product_details $\rightarrow$ exchange_delivered_order_items

Plan B: find_user_id_by_name_zip → get_user_details → get_order_details →
return_delivered_order_items

Seeding (free rollouts): after the rewritten request, the transcript contains an assistant turn “My plan: *[plan]*. I will execute it now.” followed by a user turn “Go ahead and execute your plan step by step.”

An override, verbatim (rollout B1: plan B, the return path, seeded as the agent’s own): after identifying the user and order, the agent inspected the product and two candidate items, then called

```
exchange_delivered_order_items(order_id="#W9502127",  
item_ids=["6259501109"], new_item_ids=["7958300294"],  
payment_method_id="paypal_2433177")
```

It never called the return tool named in its stated plan and instead executed the fully specified bagless-canister exchange. No verbal signal marks the substitution. All four seeded rollouts on this task realize outcome A, and all four unseeded rollouts also choose A: the override direction is exactly the task’s revealed preference.