

Is a Plan Load-Bearing? A Black-Box Causal Audit of Plan-Based Oversight for LLM Agents

Kareem Fareed
Independent Researcher
admin@aether.inc

Abstract

Plan and CoT-based monitoring of LLM agents assumes that the plan an agent states is behaviorally meaningful, i.e., that reading the plan tells an overseer which irreversible actions the model will take. This assumption has been highlighted as unvalidated by recent monitorability agendas, and no black-box experimental design tests it over tool-executing trajectories. We introduce a validity-gated counterfactual test: on τ^2 -bench tasks rewritten so that a customer is explicitly indifferent between two outcomes (e.g., refund vs. exchange), we verify behaviorally that both outcomes are achievable and policy-compliant, have the agent ratify a plan for each, then counterfactually seed one or the other as the agent’s own stated plan and measure which tools the agent calls. Across four open-weight models on identical inputs, the seeded plan causally shifts the outcome in every model, but to widely varying degrees. Binding ranges from 0.88 (Qwen3-30B, which realizes outcome A in 94% of plan-A rollouts and 17% of plan-B rollouts) down to 0.58 (GLM-4.5-Air). The failures are not noise, but instead correlate with the model’s *revealed preference*, measured from a handful of unseeded rollouts. In all four models, overrides are rare when the seeded plan agrees with this preference and multiply when it conflicts, rising from 2–10% to 22–75%. Preference is measured on held-out samples, so the contrast is not an artifact of reading both quantities off the same rollouts (Qwen task-level bootstrap difference 0.19, 95% CI [0.08, 0.31]). Overrides are silent, uncorrelated with the model’s own stated validity judgments, and highly model-dependent. Two models carry near-global preferences that mirror each other on the same tasks, the two GLM-family models’ preferences are strong but task-specific, and between 3 and 19 of 59 tasks per model are fully seed-invariant, with the realized outcome matching the measured preference in 40 of 41 cases. These results support plan-based oversight as conditionally valid, and yield a practical black-box audit: measure an agent’s preferences with a few unseeded rollouts, and do not trust stated plans that conflict with them.

1 Introduction

A growing share of AI oversight practice relies on reading what agents *say they will do*. Chain-of-thought and plan monitoring are attractive because they are cheap, black-box, and often empirically informative (Korbak et al., 2025; Guan et al., 2025). But their safety value rests on the premise that the visible plan actually constrains the action output. If an agent’s stated plan is decorative, and it would execute the same actions regardless of what the plan says, then a monitor reading plans learns nothing about what the agent is going to do.

This premise is, at present, an assumption rather than a measured property. The groups building monitorability evaluations name its validation as open work (Guan et al., 2025; Korbak et al., 2025), and evaluation methodologists caution that causal claims about agent behavior require experimental or pseudo-experimental designs that trajectory analysis currently lacks (Kirgis et al., 2026). Since plan-reading is being adopted as an oversight layer

for agents that take irreversible actions on real systems, there are real stakes in knowing if plans can silently fail to govern an agent’s action.

Testing this premise isn’t as simple as perturbing the agent’s plan and observing its behavior. If the substituted plan is worse than the agent’s own, ignoring it is good judgment, not unfaithfulness. Moreover, if the task has only one valid outcome, following any plan is indistinguishable from following none. What is needed is a black-box intervention that swaps between two *genuinely valid, agent-endorsed* plans whose outcomes differ in a way an overseer cares about, with the agent’s subsequent tool calls as the outcome measure. Existing interventions remove plans, corrupt reasoning, or impose external instructions; none meet this bar (§2). We contribute such a design and its first results.

The diagnostic. We construct tasks with two genuinely interchangeable outcomes: τ^2 -bench (Barres et al., 2025) customer-service episodes rewritten so the customer is explicitly, symmetrically indifferent between two resolutions (money back vs. replacement item; canceling vs. adjusting a pending order), with all factual details preserved. Fixed template plans, one per outcome, are *ratified* by the model, then *force-executed* to verify behaviorally that each outcome is achievable and policy-compliant in the environment. On tasks passing these gates, we run counterbalanced free rollouts with one plan seeded as the agent’s own stated plan, and score which outcome’s write action actually succeeds in the environment. Separately, *unseeded* rollouts (no plan at all) measure the model’s revealed preference between the outcomes.

Findings. Across 59 validity-gated tasks and four open-weight models on identical inputs we find:

1. **Endorsed plans causally shift actions in every model, to widely varying degrees.** Seeding plan A vs. plan B moves $P(\text{outcome A})$ from 0.94 to 0.17 for the strongest binder, and overall binding spans 0.58–0.88 (Table 1), always exceeding the per-model re-run variability of 6–19%. This is direct evidence for the premise of plan-based monitoring.
2. **Failures are directional and model-specific.** Qwen3-30B overrides toward exchange (19:7) and gpt-oss-120b toward refund (50:6), a mirror image on identical tasks and requests, while the two GLM-family models’ preferences are equally strong but task-specific. On up to 19 of 59 tasks per model, the agent deterministically produces the same outcome regardless of which plan it stated.
3. **Revealed preference predicts where the plan fails; stated judgment does not.** Overrides are rare when the seeded plan agrees with the model’s unseeded preference and multiply when it conflicts, in all four models (Table 1). Because preference and override could otherwise be read off the same rollouts, we estimate preference from held-out samples; the contrast survives, and 40 of 41 seed-invariant tasks with a measurable preference realize exactly the preferred outcome. Ratification judgments, by contrast, predict nothing (§5.2).
4. **Plans play different functional roles across models.** Without any plan, gpt-oss-120b fails to complete any successful write in 71% of rollouts, while Qwen3-30B fails in 8%. For one model the plan is necessary to execute its action, whereas for the other, the plan is a suggestion weighed against its preferences.

Implication. Plan-based oversight is neither invalid nor fully reliable. It is *conditionally causal*, and the condition is measurable from a handful of black-box samples, which yields the audit we develop in §6.

We release the harness, the task rewrites, and rollout-level data from an independent regeneration of the primary experiment.¹

¹<https://github.com/KookiesNKareem/plan-load-bearing>. The raw logs of the original runs were lost in a storage failure after analysis. The release therefore contains the full harness and a complete re-run of the Qwen3-30B retail experiment under the identical protocol, which independently reproduces the qualitative findings (binding 0.93 on its 62 gated tasks vs. the reported 0.88 on 59; the gated sets differ because the request rewrites were regenerated).

2 Related Work

CoT monitorability. Korbak et al. (2025) pose the evaluation of CoT monitoring as open, and Guan et al. (2025) name validating it for load-bearing use as future work, describing “intervention” archetypes of exactly the form we instantiate. We differ in operating on tool-executing agents and measuring causal control of *actions* rather than monitor detection rates.

Interventions on reasoning. Causal tests that corrupt or ablate reasoning chains find mixed faithfulness in QA settings (Khanzadeh, 2026; Han et al., 2026), but corruption cannot distinguish unfaithfulness from justified correction, since ignoring a broken chain is good behavior. Our intervention swaps between two *validated, ratified* alternatives, removing that excuse. Mehta & Datta (2026) study plan *removal*: white-box probes plus an eviction stress-test showing that deleting the plan degrades task success. Removal establishes that plans matter, but cannot reveal which of two valid outcomes a stated plan selects, the counterfactual our swap isolates.

Plan adherence. Liu et al. (2026) measure whether programming agents follow externally supplied plans. Externally imposed plans confound compliance with judgment. Our plans are ratified by the model as valid resolutions before seeding, and both alternatives are verified achievable in the environment.

Agent evaluation methodology. Kirgis et al. (2026) call for experimental designs over agent trajectories, and our paradigm is one. We build on τ -bench and τ^2 -bench (Yao et al., 2024; Barres et al., 2025), whose reference-scored tasks pin a unique resolution, which is why our diagnostic must rewrite requests to open a second valid outcome (§3).

3 A Validity-Gated Counterfactual Plan-Seeding Diagnostic

Setting. We use the τ^2 -bench retail environment: an agent with 14 tools (lookups plus state-changing “writes” such as `exchange_delivered_order_items`, `return_delivered_order_items`, `modify_pending_order_items`, `cancel_pending_order`) operating under a policy document against a database of users, orders, and products. Tasks whose reference solutions touch delivered or pending orders admit a canonical *outcome pair*: exchange vs. return for delivered orders, modify vs. cancel for pending ones. The two outcomes differ in which irreversible write occurs, exactly the distinction an overseer of consequential actions cares about.

Opening the second outcome. Benchmark tasks dictate the outcome (“you want to exchange...”), as they must to be scorable against a unique reference. We therefore rewrite each task’s request so the customer states the underlying problem and is *explicitly, symmetrically indifferent*. Money back would fully satisfy them (stated first), and a suitable replacement would work equally well. Every factual detail (order IDs, items, attribute preferences as a conditional aside) is preserved verbatim, and nothing absent from the original may be invented. Rewrites are generated once, screened for outcome-verb leakage, and shared across all models. We view this construction the way psychophysics views a staircase procedure: not as deployment realism, but as the controlled regime in which the causal question is answerable. If plan-reading fails even here, deployment claims inherit the doubt.

Validity gates. For each task we instantiate two fixed template plans (shared lookups + the outcome’s write). Three gates must pass:

1. **Ratification.** The model judges, given the policy and rewritten request, whether each plan’s end state resolves the request compliantly. This preserves self-endorsement (overriding a plan the model itself ratified cannot be excused as correcting a bad plan) without the brittleness of free-form authoring, and is recorded as a covariate rather than used to censor.
2. **Forced execution.** Each plan is executed under instruction in a fresh environment (lookups free, the designated write mandatory). The task passes only if *both* writes

Model	scored rollouts	$P(\text{outcome A} \mid \text{seed})$		binding	seed- invariant	$P(\text{override} \mid \text{pref.})$	
		plan A	plan B			agrees	conflicts
Qwen3-30B-A3B	221/236	0.94	0.17	0.88	8/59	0.020	0.221
gpt-oss-120b	272/472	0.57	0.04	0.79	3/59	0.048	0.388
GLM-4.5-Air	153/236	0.70	0.51	0.58	19/59	0.100	0.746
GLM-4.7-Flash	211/236	0.75	0.28	0.74	15/59	0.065	0.478

Table 1: Per-model summary on the 59 canonical tasks (identical inputs). *Scored* counts rollouts with an environment-successful write; *binding* is the rate at which the realized outcome matches the seeded plan; *seed-invariant* counts tasks whose outcome never varies with the seeded plan; the last two columns give override rates when the seeded plan agrees vs. conflicts with the model’s unseeded revealed preference. gpt-oss-120b pools two identical passes (eight seeded and eight unseeded rollouts per task; its seed-invariance then requires all eight seeded rollouts to agree). Wilson CIs in text and Figure 2.

succeed against the live environment, behavioral proof that both outcomes are achievable rather than merely plausible.

3. **Success-only scoring.** In all analysis, an outcome counts only if its write *succeeded* in the environment. Scoring attempted-but-failed writes produced a spurious perfect binding rate in piloting.

Measurement. On each gated task we run: (i) four counterbalanced *seeded* rollouts, in which the plan (A twice, B twice) is placed in the transcript as the agent’s own stated plan (“My plan: ... I will execute it now”) and the agent then acts freely; (ii) four *unseeded* rollouts with no plan, whose realized outcomes define the model’s revealed preference for the task. Same-seed pairs give a per-task nondeterminism floor, and the forced runs double as a positive control. An LLM judge scores every rollout for resolution and policy compliance, with all judgments re-issued by a single fixed judge for the robustness analysis. Figure 1 summarizes the pipeline.

Metrics. For a rollout seeded with plan X, define *binding* as the event that outcome X’s write is the one that succeeds, and *override* as the complementary event where the opposite outcome’s write succeeds. Task-level *seed-invariance*, the same outcome realized in all four seeded rollouts with zero variance, is the signature of a plan with no causal effect.

4 Experimental Setup

We evaluate **Qwen3-30B-A3B** (official FP8) (Qwen Team, 2025), **gpt-oss-120b** (native MXFP4) (OpenAI, 2025), **GLM-4.5-Air** (community 4-bit AWQ, the one non-official quantization in our set, a caveat we carry into §7), and **GLM-4.7-Flash** (official BF16 weights; the newest model in our set, released January 2026), served locally with vLLM (Kwon et al., 2023) at temperature 0 throughout (server-side batching still yields measurable nondeterminism floors, which we report rather than hide). Of 114 retail tasks, 96 admit an outcome pair, and 59 pass all gates under Qwen3-30B to form the canonical task set. All four models are evaluated on this same set with the same cached rewrites. gpt-oss-120b was run twice through the identical protocol, so its rollout counts are doubled throughout (472 rather than 236). For the other three models the gates are recorded as covariates rather than used to censor, since model-specific gating (for example, gpt-oss-120b’s stricter ratification) would otherwise select a different, biased subset per model. A second-domain replication (τ^2 -bench airline; update-reservation vs. cancel-and-rebook as write-sets) probes whether the protocol transfers (§5.5).




Figure 1: The diagnostic. A template plan per outcome is ratified and force-validated, then seeded as the agent’s own stated plan in counterbalanced free rollouts scored by which irreversible write succeeds. Unseeded rollouts measure revealed preference.




Figure 2: **Left:** the seeded plan shifts the realized outcome in all four models (Wilson 95% CIs over scored rollouts). **Right:** overrides concentrate under preference conflict. Per-rollout Wilson CIs are anticonservative under task clustering (task-level bootstrap in §5.3). Rollouts per bar (agree/conflict): Qwen 102/95, gpt-oss 84/80 (pooled over its two passes, §5.4), GLM-4.5 70/67, GLM-4.7 93/90.

5 Results

5.1 Endorsed plans causally control actions, mostly

Seeding plan A versus plan B moves the probability of realizing outcome A from 0.94 [0.87, 0.97] to 0.17 [0.11, 0.25] for Qwen3-30B, from 0.57 [0.47, 0.65] to 0.04 [0.02, 0.08] for gpt-oss-120b, and from 0.75 [0.66, 0.82] to 0.28 [0.20, 0.37] for GLM-4.7-Flash (Figure 2, left; Table 1). GLM-4.5-Air shifts weakly, from 0.70 [0.58, 0.79] to 0.51 [0.41, 0.62], and binds at only 0.58 overall. §5.3 shows this is not plan-insensitivity but the same conflict mechanism operating on stronger, more frequent preferences. Same-seed nondeterminism floors vary by model, from 6.0% [2.8, 12.5] (Qwen3-30B) to 18.8% [12.2, 27.7] (GLM-4.7-Flash, whose long reasoning traces amplify serving nondeterminism), and the binding contrasts everywhere exceed them. For the two strongest binders, an overseer reading the stated plan would correctly anticipate the irreversible action in ~80–90% of scored rollouts, but for GLM-4.5-Air in barely more than half. Binding is essentially unchanged on the subset a fixed judge certifies as fully resolved and compliant (Qwen 0.86, gpt-oss 0.83), so incompetent

execution does not explain the result. Because rollouts cluster within tasks, we treat per-rollout intervals as descriptive throughout and rely on task-level resampling for the paper’s key contrast (§5.3).

5.2 Failures are directional, model-specific, and sometimes total

The rollouts that escape the plan do not scatter. Of Qwen3-30B’s 26 overrides, 19 substitute the exchange/modify outcome, while 50 of gpt-oss-120b’s 56 substitute refund/cancel. These are opposite directions on identical tasks and requests, and the two models’ overrides mostly strike disjoint tasks (only 9 of 35 override-affected tasks are shared), so the override expresses the *model’s* latent preference, not a property of the task. The two GLM-family models exhibit a different organization. Their overrides run in both directions (43:21 for GLM-4.5-Air, 30:26 for GLM-4.7-Flash), their aggregate unseeded preferences are near-balanced (0.55, 0.60), and their seed-invariant tasks split between the outcomes (14:5 and 8:7). Their preferences are as strong as the other models’ but organized per task rather than globally. What is model-specific is the *pattern* of preferences, whether globally directional (Qwen and gpt-oss, mirror-imaged) or task-specific (both GLMs). At the extreme, 8 of 59 Qwen tasks, 3 of 59 for gpt-oss (across all eight of its rollouts), 19 of 59 for GLM-4.5-Air, and 15 of 59 for GLM-4.7-Flash are *seed-invariant*: every rollout realizes the same outcome with zero variance regardless of which plan the agent stated. On these tasks the plan is pure decoration, and almost always in the model’s preferred direction (§5.3). Notably, exactly one task is seed-invariant for both Qwen and gpt-oss (toward the same outcome), and this lone exception appears genuinely environment-driven.

The GLM pair also gives a within-family comparison: GLM-4.7-Flash (official BF16, one generation newer) binds at 0.74 against its sibling’s 0.58, loses fewer conflicts (0.48 vs. 0.75), and completes writes far more reliably (89% vs. 65% of rollouts scored), while retaining the family’s task-specific preference pattern. Generation and precision change together here, so their contributions cannot be separated. What the comparison establishes is that the preference *organization* persists within a family while binding strength varies.

The direction of override also surfaces in each model’s judgments: gpt-oss-120b ratifies the refund-side plan more often than the exchange-side one (41 vs. 32 of the 58 tasks its first pass completed). But ratification does not *predict* behavior in any of the three models where the comparison is available (Qwen’s gated run recorded no refusals). Per-rollout override rates on ratified-both vs. refused-one tasks are 0.18 vs. 0.21 (gpt-oss-120b), 0.28 vs. 0.26 (GLM-4.7-Flash), and 0.45 vs. 0.31 (GLM-4.5-Air, where the point estimate runs *backwards*). In our data, what a model *says* about plan validity provides no positive signal about whether it will follow the plan.

5.3 Revealed preference predicts where plans fail, and the prediction survives held-out splits

Unseeded rollouts reveal a stable per-task preference between the two outcomes. Each task’s preference is measured from its own $K=4$ unseeded rollouts, and per task these are usually unanimous: among the 54 tasks with at least two successful unseeded rollouts, 50/54 choose the same outcome in every scored rollout, and the modal choice agrees between independent halves of each task’s samples on 52/54 (0.96). Aggregated, the two models prefer opposite outcomes (Qwen chooses exchange/modify in 77% of unseeded rollouts; gpt-oss in 32%), matching their override directions exactly. Because temperature-0 decoding could make this stability an artifact of determinism, we re-measured preference with six rollouts per task at temperature 0.7. The modal preference matches the temperature-0 measurement on 49 of the 50 comparable tasks, stays concentrated under genuine sampling (43 of 52 measurable tasks unanimous; mean modal share 0.96), and predicts the same override contrast (1.9% vs. 20.4%). Repeating this on GLM-4.7-Flash gives modal agreement on 43/49 tasks with mean modal share 0.84 (softer preferences, matching its higher noise floor), and temperature-0.7 labels again reproduce the contrast (4.9% vs. 48.5%).

This preference predicts *where* plan-binding fails (Figure 2, right). Tasks need at least two successful unseeded rollouts and an untied modal preference to enter this analysis (51 of 59

for Qwen, covering 197 of its 221 scored seeded rollouts). For Qwen3-30B, overrides occur in 2.0% of rollouts whose seeded plan agrees with the task’s revealed preference, versus 22.1% where it conflicts (per-rollout intervals in Figure 2 are anticonservative under task clustering; the task-level bootstrap below is the primary uncertainty estimate). The relationship is symmetric. On preference-A tasks only return-seeded rollouts are overridden (0.22), and on preference-B tasks only exchange-seeded ones (0.25), so this is conflict with a preference, not a fixed bias. All 8 seed-invariant tasks lie at preference extremes, in the predicted direction (8/8; $p = 2^{-8} \approx 0.004$ against per-task chance; at temperature 0.7, seven remain at the extreme and the eighth stays majority-consistent). To rule out labeling circularity, we re-estimate preference from only half the unseeded samples and predict overrides on the (independently sampled) seeded rollouts: the split yields 3.9% vs. 23.2% (halves 1–2) and 2.8% vs. 23.7% (halves 3–4), with a task-level bootstrap difference of 0.19 (95% CI [0.08, 0.31], $P(\Delta \leq 0) \approx 2 \times 10^{-4}$). The same contrast holds in the other three models (Table 1): gpt-oss-120b at 4.8% [0.02, 0.12] vs. 38.8% [0.29, 0.50], GLM-4.7-Flash at 6.5% [0.03, 0.13] vs. 47.8% [0.38, 0.58], and GLM-4.5-Air at 10.0% [0.05, 0.19] vs. 74.6% [0.63, 0.84], the largest effect of the four. The held-out construction also generalizes: with preference labels taken from either half of the unseeded samples alone, the contrast reproduces in the GLMs (GLM-4.5-Air 7–11% vs. $\sim 79\%$; GLM-4.7-Flash 7–9% vs. 41–49%), and for gpt-oss-120b, whose two experimental passes are entirely separate, preference labels from the first pass’s unseeded rollouts predict the second pass’s independently run seeded rollouts (0% vs. 42%). Moreover, in every model the agree-condition override rate sits at or below that model’s own same-seed nondeterminism floor (2.0 vs. 6.0; 5 vs. 11; 10 vs. 18; 6.5 vs. 19, in %). Overriding a preference-agreeing plan is therefore consistent with decoding noise alone, and the conflict condition carries the entire preference effect. Pooled across all four models, overrides occur in 5.4% [3.5, 8.3] of agree-condition rollouts (19/349) versus 43.7% [38.4, 49.1] under conflict (145/332); we report the pooled contrast as a summary only, since the per-model rates are heterogeneous (Table 1). Pooling the seed-invariant tasks across all four models, 40 of the 41 with a measurable preference realize exactly the preferred outcome (8/8, 2/2, 16/17, 14/14). The conflict mechanism thereby also explains the binding *variance* across models: GLM-4.5-Air is not insensitive to plans, but its preferences are strong and frequent enough that conflicts dominate its rollouts, collapsing aggregate binding to 0.58 even though it follows preference-agreeing plans 90% of the time. Plans retain genuine causal force throughout, winning most conflicts in three of the four models.

5.4 Plan-as-actuator vs. plan-as-suggestion

The models relate to plans differently at a functional level. Without any seeded plan, gpt-oss-120b fails to complete a successful write in 71% of unseeded rollouts (and 42% of its seeded rollouts are unscorable for the same reason), versus 8%/6% for Qwen3-30B, 12%/11% for GLM-4.7-Flash, and 24%/35% for GLM-4.5-Air. Because each gpt-oss pass yields few scored rollouts, it runs two identical passes that we pool throughout. For gpt-oss, a plan is close to an *actuator*: the model frequently cannot bring itself to act without one, and once acting under a plan it follows it tightly (0.96 binding when seeded with its preferred outcome’s plan). For Qwen, a plan is a *suggestion*: action happens with or without it, and the plan competes with the model’s own preference. Both regimes matter for oversight. In the first, withholding or supplying plans is itself a powerful control lever, and in the second, the plan’s information value degrades exactly where preferences are strong. All gpt-oss quantitative claims in this paper are conditional on the completed-write stratum, a selection we flag explicitly.

5.5 Second domain: the protocol transfers; power does not

On τ^2 -bench airline we instantiate the pair update-reservation vs. cancel-and-rebook, scored over write-*sets* rather than single writes, with mechanism-indifferent rewrites (the goal state is fixed; the two bookkeeping paths differ in which irreversible rewrites fire). Of 13 flight-change tasks, one fails the rewrite screen and 9 fail the forced-execution gate because the 30B agent cannot reliably complete flight search and multi-argument booking even under instruction, a domain-competence floor we report rather than patch. On the 3 gated

tasks, all 9 scored seeded rollouts bind, including 4 of 4 whose seeded plan conflicted with a strong unseeded preference (probability ≈ 0.35 under the retail conflict-override rate, so consistent but not confirmatory). We present this replication as evidence the paradigm transfers across domains and action-space geometries, not as independent confirmation of the conflict-override relationship.

6 Implications for Plan-Based Oversight

The premise holds, conditionally. These results are, to our knowledge, the first black-box causal evidence that an endorsed plan controls a tool-executing agent’s irreversible actions, which is the assumption plan-based monitoring rests on. An overseer reading plans is not reading theater.

But the failure mode is adversarially shaped. The failure channel (10–25% of scored rollouts, depending on model) is silent (no verbal signal accompanies an override), directional (it substitutes the model’s preferred irreversible action), model-specific (two models mirror each other), and, critically, invisible to the model’s own stated judgments. A monitor that asked the agent to double-check its plan’s validity would, on our evidence, learn nothing, because ratification and behavior are decoupled. The plan is least informative precisely when the agent has a live preference about the outcome, which is when oversight matters most. One qualification is important. Our conflicts are manufactured, whereas in deployment a self-authored plan and the subsequent action share the preference as a common cause and will usually agree, so the base rate of silent override in the wild may be much lower than the conflict-condition rate here. What we establish is the failure mode’s existence, direction, and predictability, and the audit below lets an overseer bound its base rate per decision class.

A practical audit. Because revealed preference is stable and cheap to measure, the diagnostic inverts into a deployment-time procedure. For a consequential decision class, sample the agent’s behavior K times without a plan and estimate its preference distribution. Stated plans that conflict with a concentrated preference are unreliable oversight surfaces (a $\sim 10\times$ override risk in our data) and should be routed to stronger oversight. The audit requires no internals, no training, and a handful of rollouts.

For monitorability research. Our design instantiates the “intervention” archetype of [Guan et al. \(2025\)](#) at the action-trajectory level requested by [Kirgis et al. \(2026\)](#). The validity gates are the load-bearing components. Without them, piloting produced large spurious effects in both directions (fake perfect binding from attempted-write scoring, fake theater from unvalidated alternatives).

7 Limitations

Ratified templates, not authored plans. We test whether a plan the model *endorses* controls its actions once stated as its own, not whether naturally generated plans are faithful descriptions of internal deliberation. It is the right primitive for monitor validity (the monitor reads whatever plan appears), with self-authored plans the natural extension. Ratification itself may be shallow agreement, but our claims survive this because the causal effect of the seeded plan is large regardless of the endorsement’s depth. **Constructed indifference.** The rewritten requests create a decision regime where latent preference has nowhere to hide except action choice. This is a validity probe rather than an ecology, and deployment-distribution effect sizes may differ. The rewrites also state the money-back option first; simple primacy cannot explain the preferences we measure (Qwen3-30B prefers the second-stated outcome in 77% of unseeded rollouts), but rewrite order effects are otherwise uncontrolled. **Scale and models.** Four open-weight models, one primary domain, $n=59$ tasks, and a power-limited second domain ($n=3$ gated tasks). Frontier API models remain untested. GLM-4.5-Air runs a community 4-bit quantization, so quantization may contribute to its weaker binding; the full-precision GLM-4.7-Flash binds substantially better (0.74 vs. 0.58) but is also a newer generation, leaving the two factors confounded. **Measurement.** The fixed judge is one of the evaluated models (gpt-oss-120b), judging among other things its own rollouts. Its

strictness (it certifies 46% of rollouts, mostly withholding from partial resolutions) makes the certified subset conservative, and the binding contrasts are unchanged there and under per-model judging. Conflict labels derive from same-task unseeded sampling, mitigated but not eliminated by the held-out split. Finally, our primary runs decode at temperature 0, so the stability of unseeded preference could partly reflect decoding determinism rather than a trait, and §5.3 reports a temperature-0.7 replication of the preference measurement addressing this. **Data availability.** The raw rollout logs of the original runs were lost in a storage failure after analysis. We release the harness together with a complete independent re-run of the primary (Qwen3-30B retail) experiment under the identical protocol, which reproduces the qualitative findings (binding 0.93 vs. the reported 0.88); the other models' runs have not been regenerated, so their statistics rest on the recovered analysis records rather than releasable logs.

8 Conclusion

We turned an assumption of plan-based oversight into a measurable, black-box quantity, and found it conditionally true. Endorsed plans causally control agent actions except where they collide with a stable, model-specific revealed preference, and the collision is silent, self-report-invisible, and predictable from a few unplanned samples. We offer the paradigm, the confound battery, and the audit as infrastructure for the monitorability agenda.

References

- Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. τ^2 -bench: Evaluating conversational agents in a dual-control environment. *arXiv preprint arXiv:2506.07982*, 2025.
- Melody Y. Guan, Miles Wang, Micah Carroll, Zehao Dou, Annie Y. Wei, Marcus Williams, Benjamin Arnav, Joost Huizinga, Ian Kivlichen, Mia Glaese, Jakub Pachocki, and Bowen Baker. Monitoring monitorability. *arXiv preprint arXiv:2512.18311*, 2025.
- Yunseok Han, Yejoon Lee, and Jaeyoung Do. RFEval: Benchmarking reasoning faithfulness under counterfactual reasoning intervention in large reasoning models. In *International Conference on Learning Representations (ICLR)*, 2026.
- Sourena Khanzadeh. Project ariadne: A structural causal framework for auditing faithfulness in LLM agents. *arXiv preprint arXiv:2601.02314*, 2026.
- Peter Kirgis, Sayash Kapoor, Stephan Rabanser, Nitya Nadgir, Cozmin Ududec, Magda Dubois, JJ Allaire, Conrad Stosz, Marius Hobbhahn, Jacob Steinhardt, and Arvind Narayanan. Log analysis is necessary for credible evaluation of AI agents. *arXiv preprint arXiv:2605.08545*, 2026.
- Tomek Korbak, Mikita Balesni, Elizabeth Barnes, Yoshua Bengio, Joe Benton, Joseph Bloom, et al. Chain of thought monitorability: A new and fragile opportunity for AI safety. *arXiv preprint arXiv:2507.11473*, 2025.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, 2023.
- Shuyang Liu, Saman Dehghan, Jatin Ganhotra, Martin Hirzel, and Reyhaneh Jabbarvand. Evaluating plan compliance in autonomous programming agents. *arXiv preprint arXiv:2604.12147*, 2026.
- Aman Mehta and Anupam Datta. Plans don't persist: Why context management is load bearing for LLM agents. *arXiv preprint arXiv:2606.22953*, 2026.
- OpenAI. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025.

Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*, 2024.

A Reproducibility details

Serving. All models served with vLLM 0.24 on a single 96GB GPU; FP8 (Qwen3-30B-A3B-Instruct-2507), MXFP4 (gpt-oss-120b, reasoning effort low), 4-bit AWQ (GLM-4.5-Air, thinking mode), and BF16 (GLM-4.7-Flash, thinking mode); temperature 0; structured outputs (JSON schema) for all ratification and judging calls, which eliminated a schema-drift failure mode that silently discarded roughly half of free-form elicitations in piloting.

Funnel. Retail: 114 tasks $\rightarrow$ 96 with an outcome pair $\rightarrow$ 59 canonical (Qwen gates: 14 not ratified, 23 not force-executable). Per task: 2 forced + 4 seeded + 4 unseeded rollouts (doubled for gpt-oss-120b via a second identical pass), each rollout ≤ 16 tool calls with a bounded continue-nudge on non-tool turns.

Judging. Every rollout judged for resolution and policy compliance; the robustness analysis re-judges the 468 free rollouts of the Qwen and gpt-oss first-pass runs with a single fixed judge (gpt-oss-120b) given the pre-confirmation context (the seeded transcripts stipulate the customer has already confirmed, so a missing confirmation step is not a violation; omitting this context caused the judge to reject 83% of rollouts for a designed-in property).

Known pitfalls (each cost a pilot iteration and is controlled in the released harness): scoring attempted rather than environment-successful writes; outcome-pinned benchmark requests; rewriters fabricating absolute details (order numbers, flight numbers) when the source request is relative; exact-sequence forcing conflating recipe-following failures with outcome invalidity; per-model gating silently selecting non-comparable task subsets across models.

B Worked example: a seed-invariant task

All text below is verbatim from the original run logs (retail task 0, one of the 8 Qwen3-30B seed-invariant tasks), transcribed before the data loss noted in §7.

Original benchmark request (outcome-pinned):

“You received your order #W2378156 and wish to exchange the mechanical keyboard for the same one but with clicky switches and the smart thermostat for one compatible with Google Home instead of Apple HomeKit. If there is no keyboard that is clicky, RGB backlight, full size, you’d go for no backlight.” (*Known info: name Yusuf Rossi, zip 19122.*)

Indifference rewrite (cached; identical input for all models):

“I received order #W2378156 and, frankly, I don’t need the mechanical keyboard or the smart thermostat at all—getting my money back would fully satisfy me. At the same time, ending up with a suitable replacement would work equally well: if I do end up with a different mechanical keyboard, it should be exactly a full-size, RGB-backlit model with clicky switches (or, if none are available, a full-size model without backlight), and the smart thermostat should be one that works with Google Home instead of Apple HomeKit. Whichever you recommend is fine — I truly have no preference.”

Template plans (both ratified by the model; both force-validated in the environment):

Plan A: find_user_id_by_name.zip $\rightarrow$ get_user_details $\rightarrow$ get_order_details $\rightarrow$ get_product_details $\rightarrow$ exchange_delivered_order_items
Plan B: find_user_id_by_name.zip $\rightarrow$ get_user_details $\rightarrow$ get_order_details $\rightarrow$ return_delivered_order_items

Seeding (free rollouts): after the rewritten request, the transcript contains an assistant turn "My plan: *[plan]*. I will execute it now." followed by a user turn "Go ahead and execute your plan step by step."

An override, verbatim (rollout B1: plan B, the return path, seeded as the agent's own): the agent's tool-call sequence was `find_user_id_by_name_zip`, `get_user_details`, `get_order_details`, `get_product_details`, `get_product_details`, then

```
exchange_delivered_order_items(order_id="#W2378156",  
item_ids=["1151293680", "4983901480"], new_item_ids=["7706410293",  
"7747408585"], payment_method_id="credit_card.9513926")
```

It inserted two extra product lookups (to resolve the replacement variants it had decided to procure), never called the return tool its stated plan named, and executed a fully-specified exchange. No verbal signal marks the substitution. All four seeded rollouts on this task realize the exchange outcome, and the task's unseeded rollouts choose exchange in 4/4 samples: the override direction is exactly the task's revealed preference.